

"Misread Spec"

The model misreads what the requirement asks, or points to something that satisfies a different requirement.

I.e., pointing to authentication for something about data integrity

"Added Requirement",

The model adds controls or hallucinates those which the requirement is not about, making up new capabilities that the requirement never mentions.

"Overthink Edge",

The model fixates on more detail or a rare edge of something the requirement does cover, when the ground truth shows it is OK to skip it.

Write this better, and mention that something can be part of the requirement, but based on the ground truth it is OK not to have it, yet the LLM emphasizes it. For example, for safeguarding passwords, the ground truth mentions the password is hashed, so its ok, but the LLM says the password does not use PBKDF, so it's not satisfied. PBKDF obviously helps the guardian, and it is fine to use it, but the LLM focuses too much on this specific case, which the ground truth doesn't emphasize.

"Assumed Type",

"Imagined Runtime",

"Performance Nitpick",

"Readability Nitpick",

"Precision Error",

"Boundary Error",

"Logic Error",

The model is wrong about the code, it misses something that is actually there or misstates a capability.

"Vague Description",

The model rejects a vague concern with no concrete example.

I.e., downgrading for “some legal phrasing” without naming any unclear wording

"Other",